

巨量資料於 R軟體中的處理與分析

(High-Performance Computing with R)

吳漢銘

國立政治大學 統計學系



<https://hmwu.idv.tw>



- Why Are R Programs Sometimes Slow?
- Profiling: Measuring R Code's Performance
- (I) Speed Up R
 - (A) Simple Tweaks
 - (B) Using Compiled Code for Greater Speed
 - (C) Using GPUs to Run R Even Faster
- (II) Use Less RAM
 - (A) Simple Tweaks
 - (B) Processing Large Datasets with Limited RAM
- (III) Multiplying Performance with Parallel Computing
- (IV) Offloading Data Processing to Database Systems
- (V) Other Consideration
- (VI) R and Big Data: Hadoop Techniques
- Conclusion

The Definition of Big Data is a Moving Target

3/66

<http://www.opentracker.net/article/definitions-big-data>

opentracker
event tracking and analytics

PRODUCTS FEATURES PRICING

Start your free, **no-risk**, 4 week trial

Definitions of Big Data

Q: Can you please provide me with a definition of Big Data?

A: The definition of Big Data is a moving target.

In order to make it possible to follow the discussion, as it evolves, we have started a list of definitions of Big Data, as we read them on the internet.

Author names: [Andrew Brust](#) (ZDNet), [Bill Franks](#) (FCW article), [PCMag](#), [Weather.com](#), [Mike](#), [O'Reilly](#), [Library](#), [Rous](#), [O'Reilly](#), [Slash](#), [O'Reilly](#), [Evels](#), [Steph](#)

31. "Big data is a popular term used to describe the exponential growth of information, both structured and unstructured. Ultimately, regardless of the factors involved, we believe that Big Data applies (per Gartner's assessment) whenever an organization's data exceeds its current capacity." [SAS](#).

32. The simplest definition of "Big Data" is "it doesn't fit in Excel" - from the full quote; "I have joked that the simplest definition of Big Data is - and when you think of it, it's true for most people who have tried the traditional approach to a Big Data one." [Stephane Hamel](#) [comment 8/2012](#) Big Data - What It Means

33. More to follow...

updated: April 26, 2013

Application Delivery Strategies

Date: 6 February 2001

File: 949
Author: Doug Laney

3D Data Management: Controlling Data Volume, Velocity, and Variety. Current business conditions and mediums are pushing traditional data management principles to their limits, giving rise to novel, more formalized approaches.

Report | [McKinsey Global Institute](#)

Big data: The next frontier for innovation, competition, and productivity

May 2011 | by James Manyika, Michael Chui, Brad Brown, Jacques Bughin, Richard Dobbs, Charles Roxburgh, Angela Hung Byers

Download [Executive Summary](#) [Full Report](#) [Kindle](#) [eBook](#)
PDF-922KB PDF-6MB MOBI-4MB EPUB-3MB

The amount of data in our world has been exploding, and analyzing large data sets—so-called big data—will become a key basis of competition, underpinning new waves of productivity growth, innovation, and consumer surplus, according to research by MGI and McKinsey's Business Technology Office. Leaders in every sector will have to grapple with

[PDF](#) [Print](#) [E-mail](#) [Share](#)

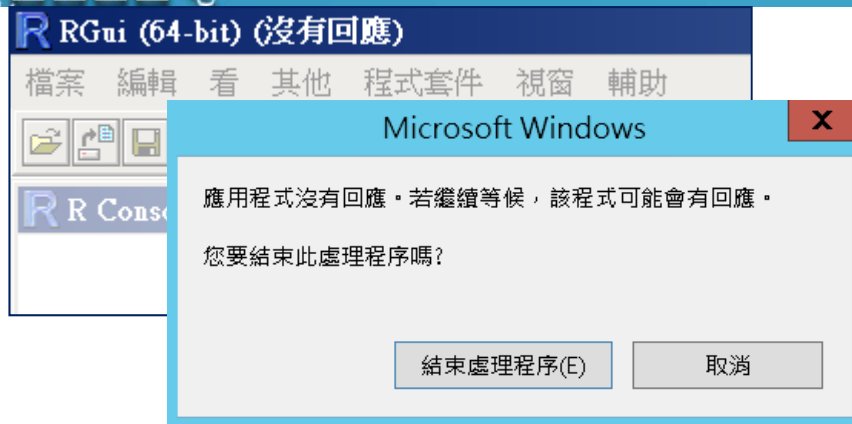
[McKinsey Global Institute](#)

[McKinsey&Company](#)

For data that is larger than the machine memory, we consider it as Big Data.


實際狀況 in R

4/66



```
R Console
> n <- 1e7
> p <- 2000
> myData <- as.data.frame(matrix(rnorm(n*p), ncol = p, nrow=n))
Error: cannot allocate vector of size 149.0 Gb
In addition: Warning messages:
1: In rnorm(n * p) :
  Reached total allocation of 32722Mb: see help(memory.size)
```

- What gets more difficult when data is big?
 - The data may not load into memory
 - analyzing the data may take a long time
 - visualizations get messy
 - etc.






High-Performance and Parallel Computing with R

<https://cran.r-project.org/web/views/HighPerformanceComputing.html>

- Parallel computing: Explicit parallelism, Implicit parallelism, Grid computing, Hadoop, Random numbers, Resource managers and batch schedulers, Applications, GPUs.
- Large memory and out-of-memory data
- Easier interfaces for Compiled code
- Profiling tools

```
> version
platform      x86_64-w64-mingw32
arch          x86_64
os            mingw32
system        x86_64, mingw32
status
major         3
minor         2.2
year          2015
month         08
day           14
svn rev       69053
language      R
version.string R version 3.2.2 (2015-08-14)
nickname      Fire Safety
> Sys.getenv("R_ARCH")
[1] "/x64"
```



Why Are R Programs Sometimes Slow?

6/66

- Three constraints on computing performance: **CPU**, **RAM**, and **disk I/O**.
- R is interpreted on the fly.
- R is single-threaded.
- R requires all data to be loaded into memory.
- Algorithm design affects time and space complexity.

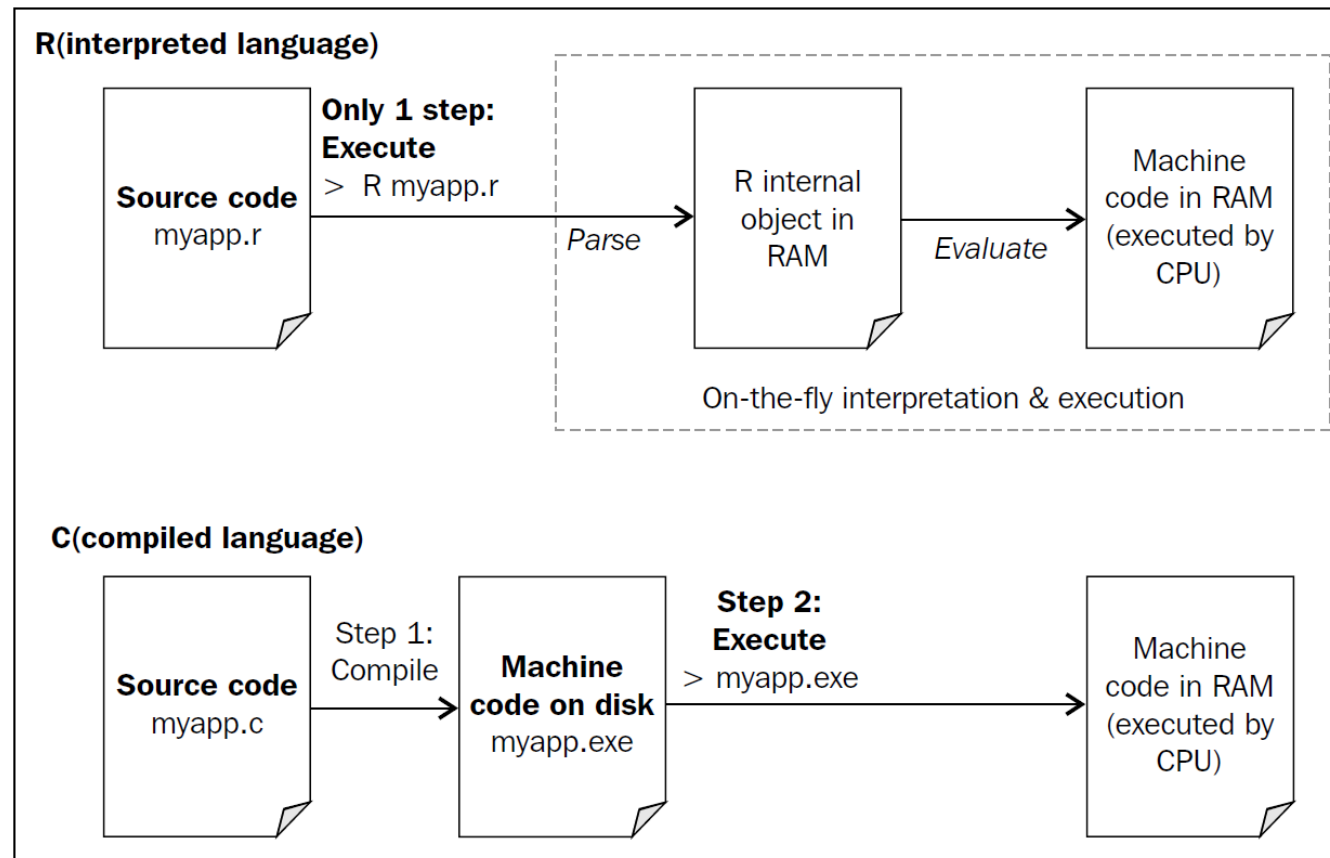


Source: <https://www.dreamstime.com>

R is Interpreted On the Fly

7/66

- R code runs relatively slow because it is reinterpreted every time you run it, even when it has not changed.
- For compiled language, once the code has been compiled, it runs very quickly on the CPU since it is already in the computer's native language.



Source: Aloysius Lim, and William Tjhi, R High Performance Programming, Packt Publishing, January 30, 2015.



Memory Allocation in R

- 當R啟動時，設定最大可獲得的記憶體：

"C:\Program Files\R\R-3.2.2\bin\x64\Rgui.exe" --max-mem-size=2040M

- ❑ 最小需求是32MB.
- ❑ R啟動後僅可設定更高值，不能再用`memory.limit`設定較低的值。

```
> # 目前使用的記憶體量
> memory.size(max = FALSE)
[1] 3845.87
>
> # 從作業系統可得到的最大量記憶體
> memory.size(max = TRUE)
[1] 3846.25
>
> # 列出目前記憶體的限制
> memory.limit(size = NA)
[1] 16343
>
> # 設定新的記憶體限制為 1024 MB
> memory.limit(size = 1024)
[1] 16343
Warning message:
In memory.limit(size = 1024) : 無法減少記憶體限制：已忽略
```

■ R與Windows作業系統

最大可獲得的記憶體

- 32-bit R + 32-bit Windows: 2GB.
- 32-bit R + 64-bit Windows: 4GB.
- 64-bit R + 64-bit Windows: 8TB.



`object.size{utils}`

- 儲存R物件所佔用的記憶體估計。

```
object.size(x)
```

```
print(object.size(x), units = "Mb")
```

```
> n <- 10000
> p <- 200
> myData <- as.data.frame(matrix(rnorm(n*p), ncol = p, nrow=n))
> print(object.size(myData), units = "Mb")
15.3 Mb

> write.table(myData, "myData.txt") ## 約 34.7 MB

> InData <- read.table("myData.txt")
> print(object.size(InData), units = "Mb")
15.6 Mb
```

NOTE: Under any circumstances, you cannot have more than $2^{31}-1=2,147,483,647$ rows or columns.




object.size{utils}

10/66

object.size (n by p, numeric)

■ KB ■ MB ■ GB



$(n * p * 8) / (1024 * 1024)$ MB

p
1 Bit = Binary Digit; 8 Bits = 1 Byte; 1024 Bytes = 1 Kilobyte; 1024 Kilobytes = 1 Megabyte
1024 Megabytes = 1 Gigabyte; 1024 Gigabytes = 1 Terabyte; 1024 Terabytes = 1 Petabyte



Measuring execution time: `system.time{base}`

11/66

```
myFun <- function(n){  
  for(i in 1:n){  
    x <- x + i  
  }  
  x  
}
```

```
> start.time <- Sys.time()  
> ans <- myFun(10000)  
> end.time <- Sys.time()  
> end.time - start.time  
Time difference of 0.0940001 secs
```

```
> system.time({  
+   ans <- myFun(10000)  
+ })  
      user  system elapsed  
0.04    0.00    0.05
```

See also: `microbenchmark`, `rbenchmark` packages

```
myPlus <- function(n){  
  x <- 0  
  for(i in 1:n){  
    x <- x + sum(rnorm(i))  
  }  
  x  
}
```

```
myProduct <- function(n){  
  x <- 1  
  for(i in 1:n){  
    x <- x * sum(rt(i, 2))  
  }  
  x  
}
```

```
> system.time({  
+   a <- myPlus(5000)  
+ })  
      user  system elapsed  
3.87    0.00    3.91  
> system.time({  
+   b <- myProduct(5000)  
+ })  
      user  system elapsed  
10.36    0.00    10.42
```

Profiling Memory Utilization:

12/66

Rprof{utils}

```
> Rprof("Rprof-mem.out", memory.profiling = TRUE)
> ans <- myFun2(5000)
> Rprof(NULL)
> summaryRprof("Rprof-mem.out", memory = "both")
```

\$by.self

	self.time	self.pct	total.time	total.pct	mem.total
"rt"	3.60	70.87	3.60	70.87	112.6
"rnorm"	1.44	28.35	1.44	28.35	108.8
"sum"	0.04	0.79	0.04	0.79	2.0

\$by.total

	total.time	total.pct	mem.total	self.time	self.pct
"myFun2"	5.08	100.00	223.5	0.00	0.00
"myProduct"	3.62	71.26	113.2	0.00	0.00
"rt"	3.60	70.87	112.6	3.60	70.87
"myPlus"	1.46	28.74	110.3	0.00	0.00
"rnorm"	1.44	28.35	108.8	1.44	28.35
"sum"	0.04	0.79	2.0	0.04	0.79

\$sample.interval

[1] 0.02

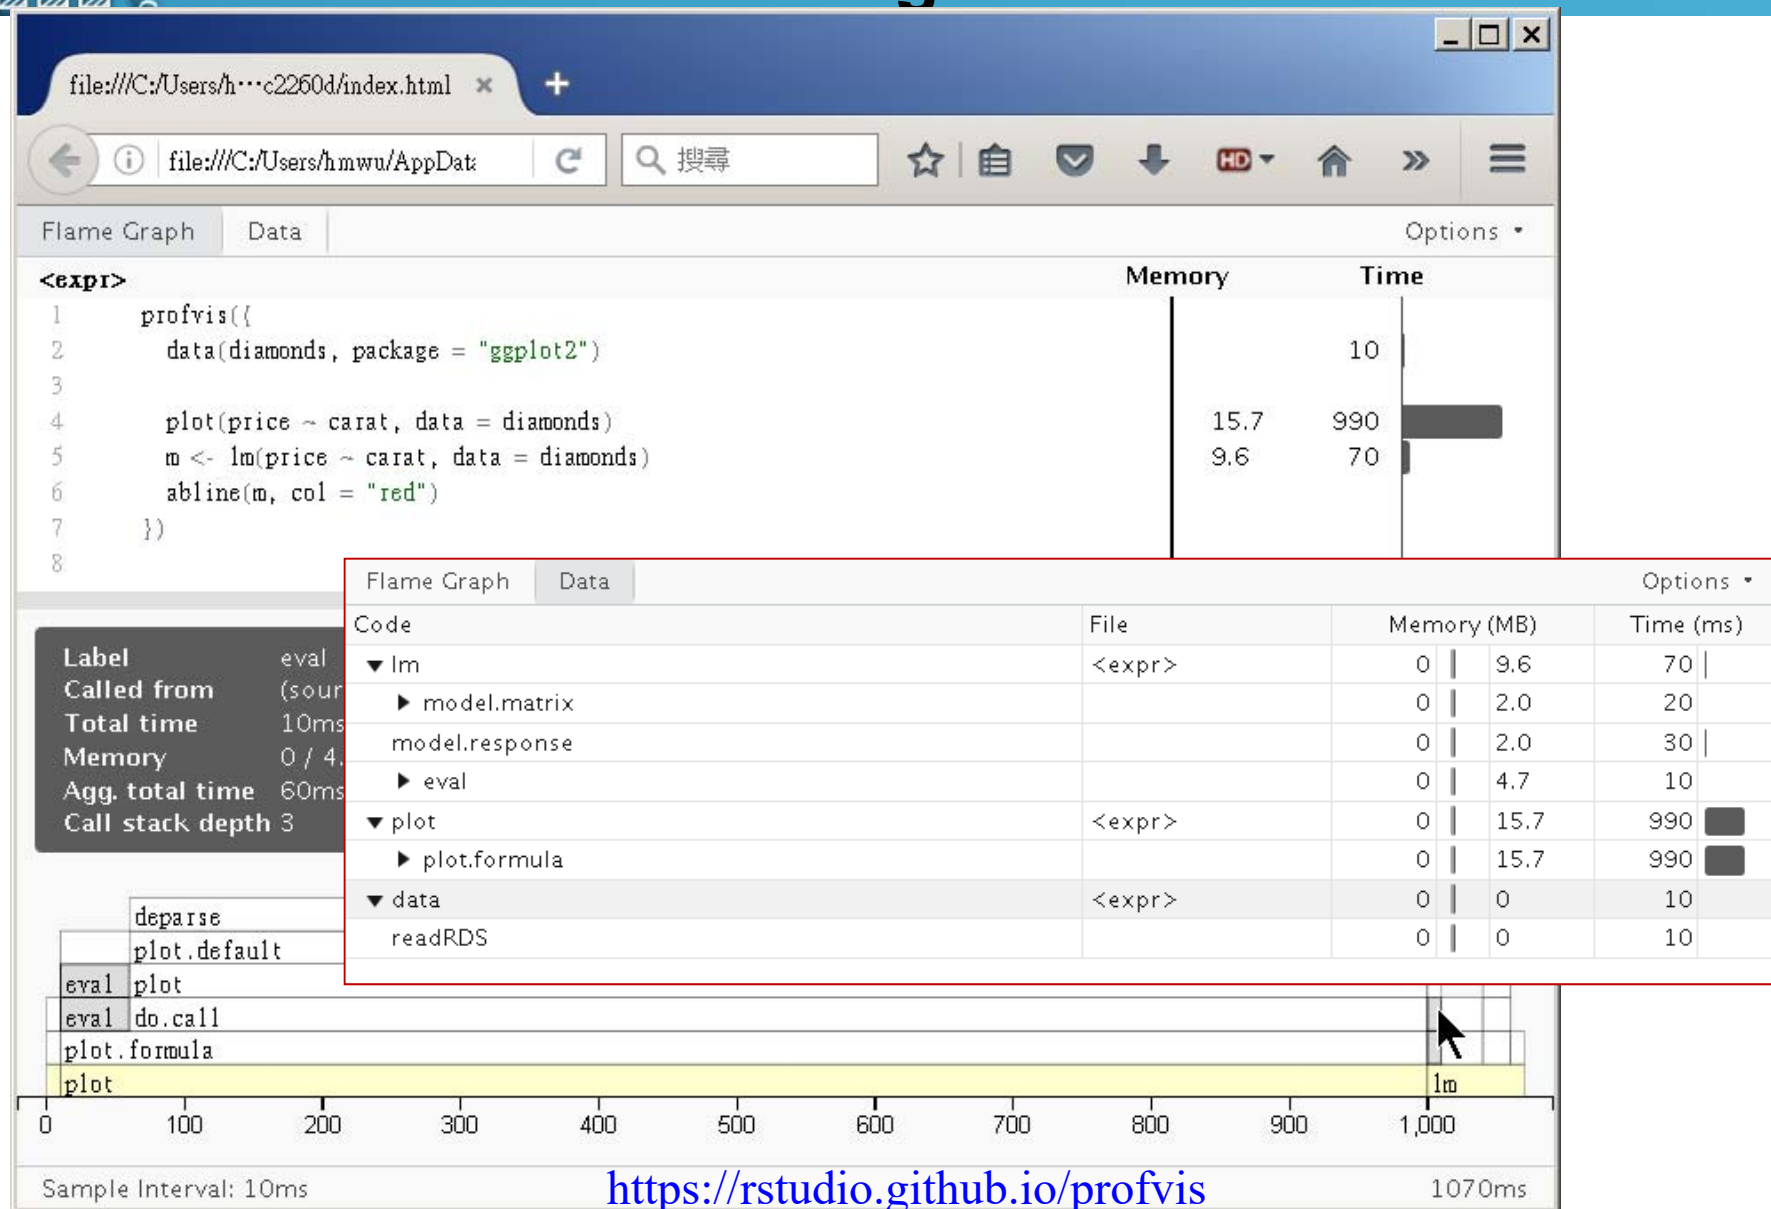
\$sampling.time

[1] 5.08

```
myFun2 <- function(n){
  a <- myPlus(n)
  b <- myProduct(n)
  list(a, b)
}
```

<See also the [proftools](#) package>

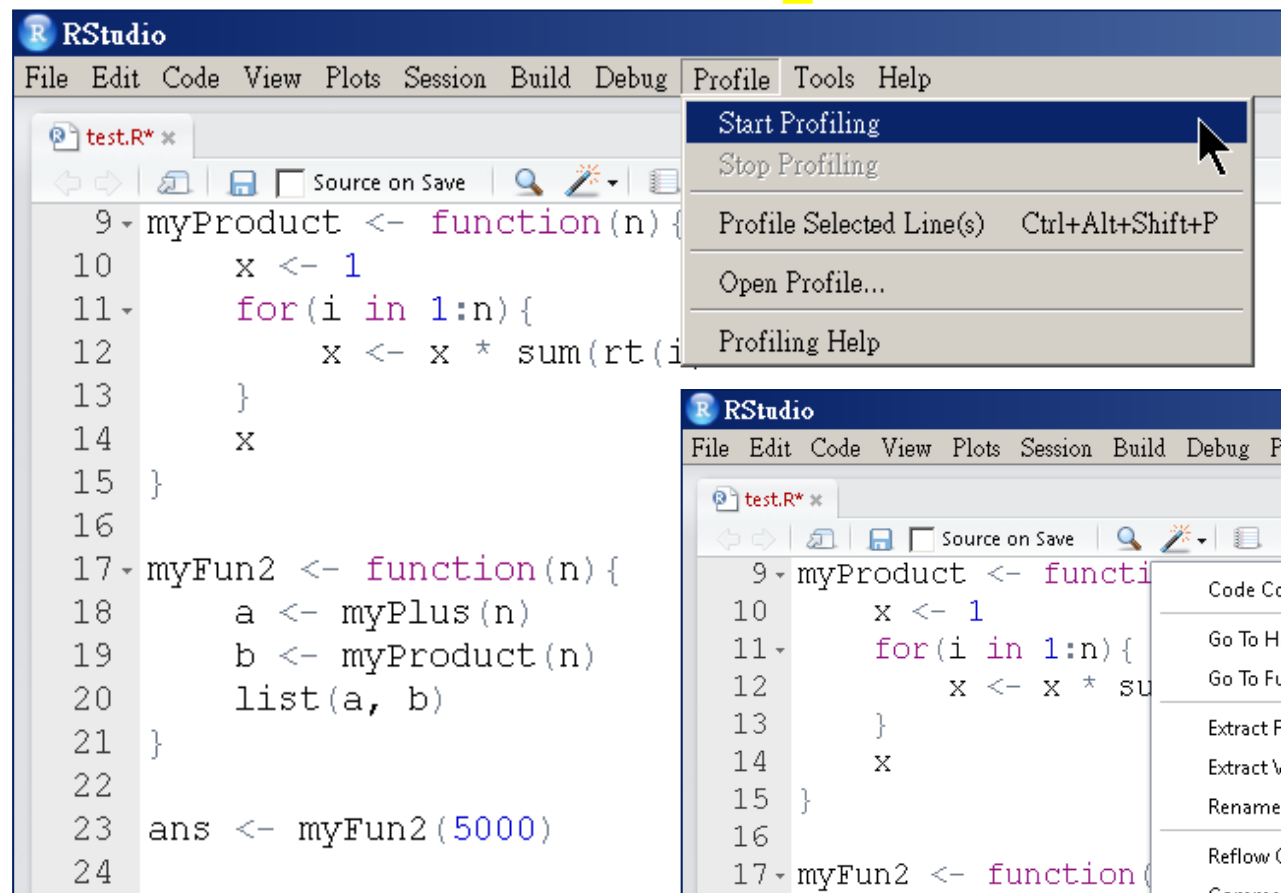
profvis: Interactive Visualizations for Profiling R Code



<https://rstudio.github.io/profvis>

Profiling with RStudio and profvis

14/66

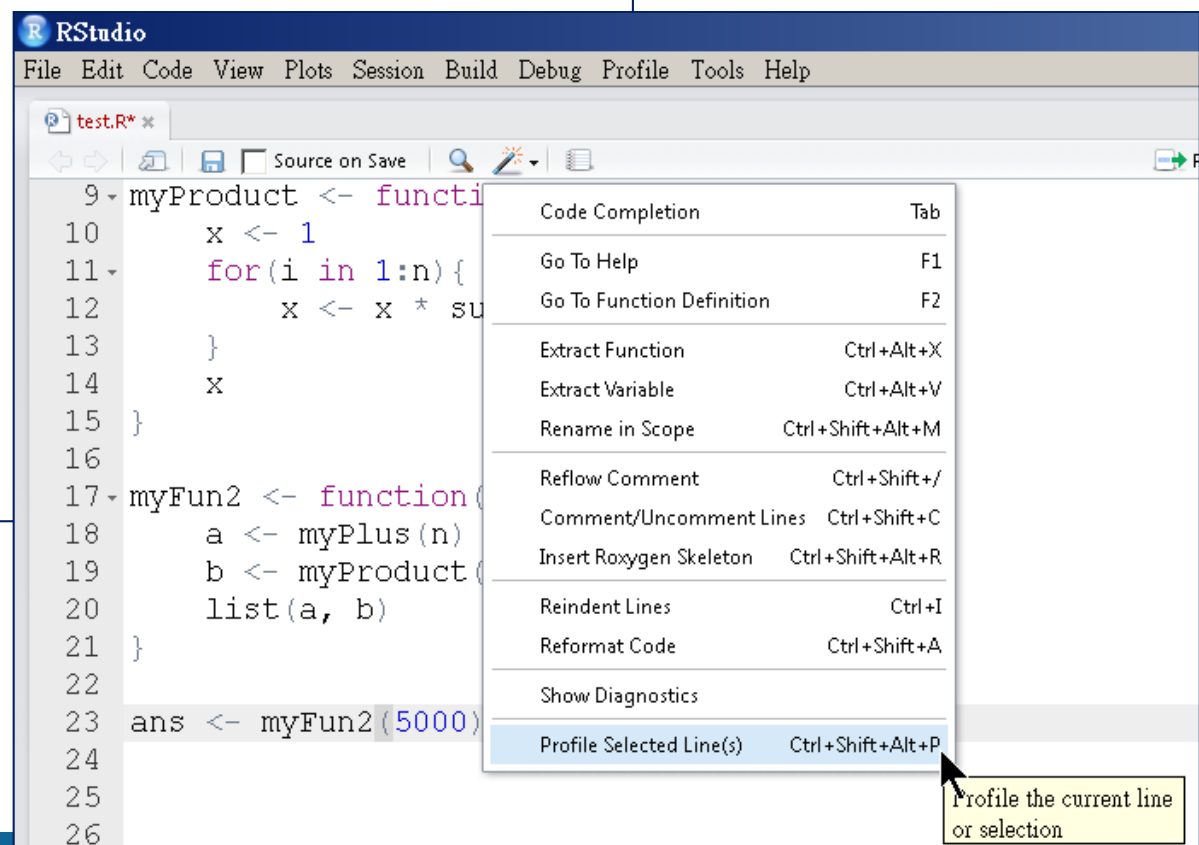


The image shows the RStudio interface with the 'Profile' menu open. The menu options are: Start Profiling, Stop Profiling, Profile Selected Line(s) (Ctrl+Alt+Shift+P), Open Profile..., and Profiling Help. The background code in the editor is as follows:

```
9 myProduct <- function(n) {  
10   x <- 1  
11   for(i in 1:n){  
12     x <- x * sum(rt(i  
13   }  
14   x  
15 }  
16  
17 myFun2 <- function(n) {  
18   a <- myPlus(n)  
19   b <- myProduct(n)  
20   list(a, b)  
21 }  
22  
23 ans <- myFun2(5000)  
24
```

RStudio v0.99.1208 Preview
(2016-06-03)

<https://rstudio.github.io/profvis/>



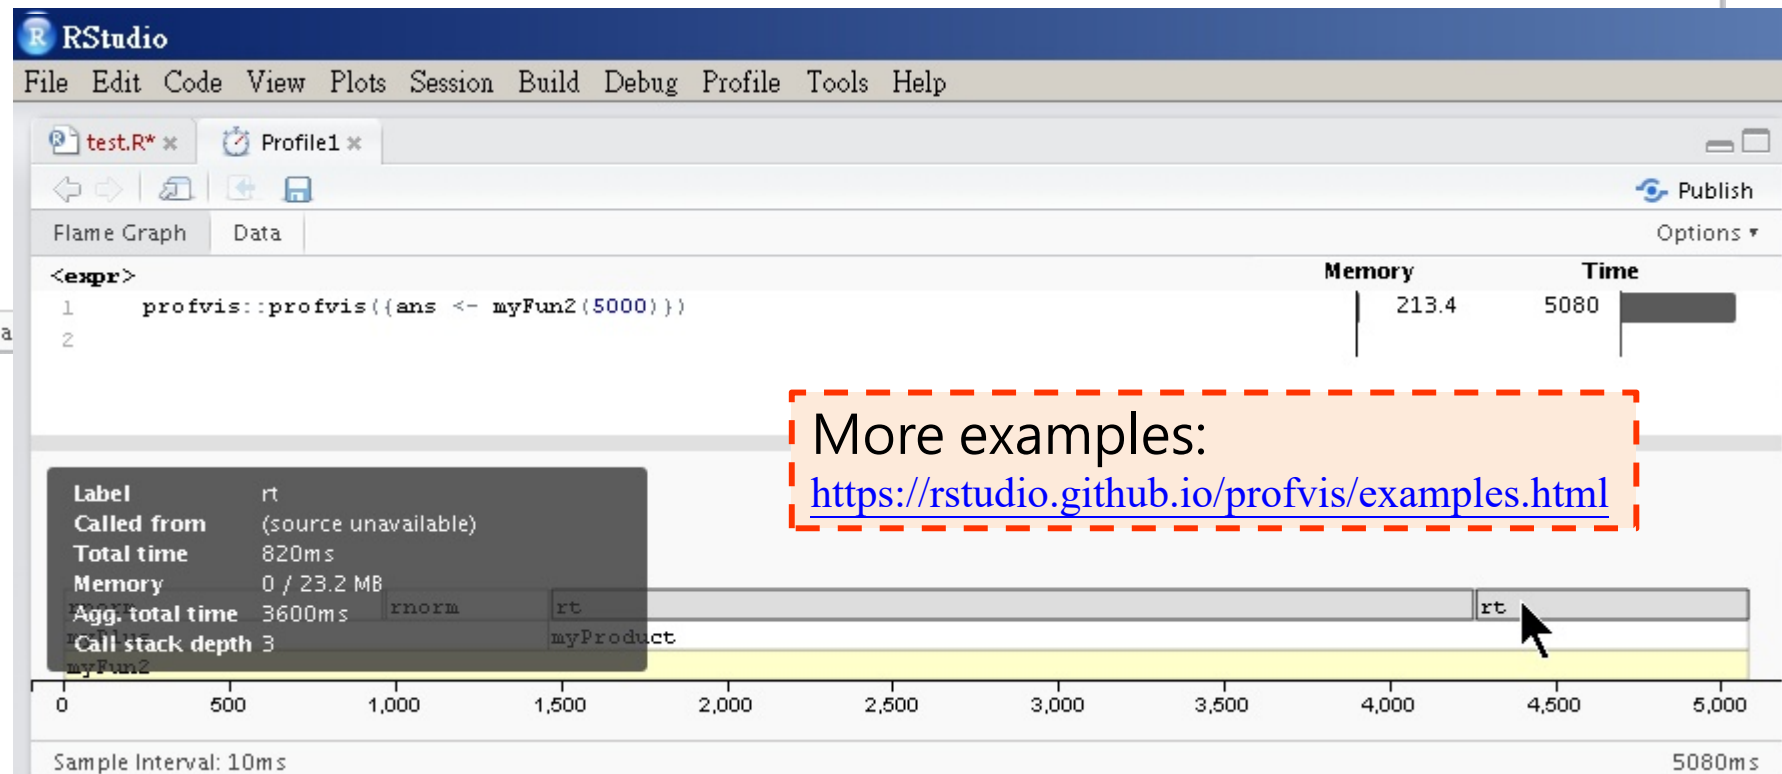
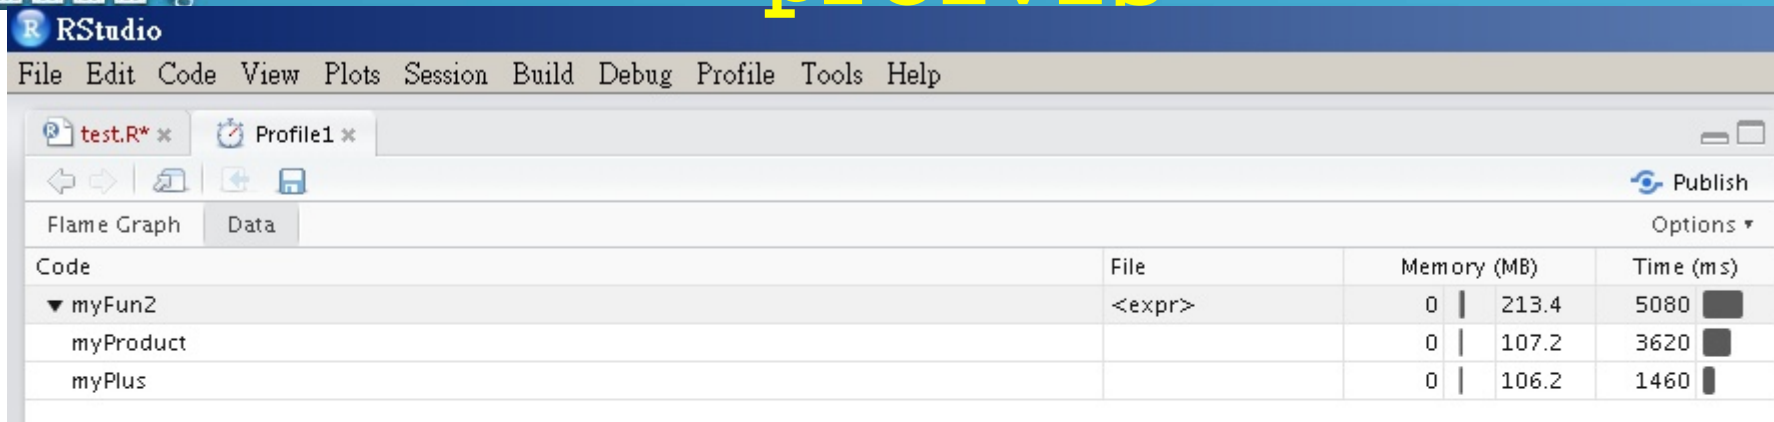
The image shows the RStudio interface with the 'Code Completion' menu open. The menu options are: Code Completion (Tab), Go To Help (F1), Go To Function Definition (F2), Extract Function (Ctrl+Alt+X), Extract Variable (Ctrl+Alt+V), Rename in Scope (Ctrl+Shift+Alt+M), Reflow Comment (Ctrl+Shift+/,), Comment/Uncomment Lines (Ctrl+Shift+C), Insert Roxygen Skeleton (Ctrl+Shift+Alt+R), Reindent Lines (Ctrl+I), Reformat Code (Ctrl+Shift+A), Show Diagnostics, and Profile Selected Line(s) (Ctrl+Shift+Alt+P). The background code in the editor is as follows:

```
9 myProduct <- functi  
10   x <- 1  
11   for(i in 1:n){  
12     x <- x * su  
13   }  
14   x  
15 }  
16  
17 myFun2 <- function(  
18   a <- myPlus(n)  
19   b <- myProduct(  
20   list(a, b)  
21 }  
22  
23 ans <- myFun2(5000)  
24  
25  
26
```

An arrow points to the 'Profile Selected Line(s)' option with the text: 'Profile the current line or selection'.

Profiling with RStudio and profvis

15/66



More examples:

<https://rstudio.github.io/profvis/examples.html>



(I) Speed Up R

(A) Simple Tweaks to Make R Run Faster

(A1) Vectorization

(A2) Use of built-in functions

(A3) Preallocating memory

(A4) Use of simpler data structures

(A5) Use of hash tables for frequent lookups on large data

(A6) Seeking fast alternative packages in CRAN

- These tweaks should be taken as the **first steps** in order to optimize an R code.
(養成良好程式設計習慣及風格)

(A1) Vectorization

17/66

- In essence, vectorization allows R operators to take vectors as arguments for quick processing of multiple values.

```
> # Non-vectorized
> n <- 1E6
> mydata <- rnorm(n)
> system.time({
+   s1 <- 0
+   for(j in 1:n){
+     s1 <- s1 + mydata[j]^2
+   }
+   s1
+ })
user system elapsed
0.97    0.00    0.97
```

```
> #Vectorized
> system.time(s2 <- sum(mydata^2))
user system elapsed
0        0        0
```

Vector size	100,000	1,000,000	10,000,000	100,000,000
Non-vectorized	120 ms	1.19 s	11.9 s	117 s
Vectorized	508 µs	5.67 ms	52.5 ms	583 ms

Source: Aloysius Lim, and William Tjhi, R High Performance Programming, Packt Publishing, January 30, 2015.



(A2) Use of Built-in Functions

18/66

- R and some CRAN packages provide a rich set of functions that are implemented in compiled languages such as C/C++.

```
> # use apply
> n <- 1E4
> p <- 1000
> mydata <- matrix(rnorm(n*p), nrow=n, ncol=p)
> system.time(data_sum1 <- apply(mydata, 1, sum))
  user  system elapsed 
0.30    0.02    0.31 
> # rowSums is an optimized and precompiled C function
> system.time(data_sum2 <- rowSums(mydata))
  user  system elapsed 
0.04    0.00    0.03
```

- Use the **optimized BLAS** (Basic Linear Algebra Subprograms) libraries (the Mac OS X version of R comes enabled with the optimized BLAS.)

```
data <- rnorm(1E7)
dim(data) <- c(1E4, 1E3)
system.time(data_mul <- t(data) %*% data)
## user system elapsed
## 7.123  0.015  7.136
system.time(data_mul <- t(data) %*% data) # with optimized BLAS
## user system elapsed
## 1.304  0.005  0.726
```

Source: Aloysius Lim, and William Tjhi, R High Performance Programming, Packt Publishing, January 30, 2015.

(A3) Preallocating Memory

19/66

- **Dynamic memory allocation** slows down a program. Every time a vector is resized, the program needs to perform extra steps that include copying the vector to a larger or smaller memory block and deleting the old vector. These steps are not needed if the memory is preallocated.

```
> n <- 1E4; dataV1 <- 1
> system.time({
+   for (j in 2:n) {
+     dataV1 <- c(dataV1, dataV1[j-1] + sample(-5:5, size=1))
+   }
+ })
user system elapsed
0.20    0.00    0.21
```

```
>
> dataV2 <- numeric(n)
> dataV2[1] <- 1
> system.time({
+   for (j in 2:n) {
+     dataV2[j] <- dataV2[j-1] + sample(-5:5, size=1)
+   }
+ })
user system elapsed
0.07    0.00    0.06
```

Vector size	10	100	1000	10,000
Dynamic allocation	0	0.006	0.288	25.373
Preallocated	0.001	0.006	0.062	0.577

Source: Lim and Tjhi, R High Performance Programming, Packt Publishing, January 30, 2015.



(A4) Use of Simpler Data Structures

20/66

- A `data.frame` object allows more flexibility than a `matrix` by allowing variables of different types.
- Applying a matrix operation on a `data.frame` is slower than on a `matrix`. One of the reasons is that most matrix operations first coerce the `data.frame` into a `matrix` before performing the computation.

```
> # working on matrix
> n <- 1E4
> p <- 1000
> mydata.mt <- matrix(rnorm(n*p), nrow=n, ncol=p)
> system.time(rs1 <- rowSums(mydata.mt))
  user  system elapsed 
 0.03    0.00    0.03 
> # working on data.frame
> mydata.df <- data.frame(mydata.mt)
> system.time(rs2 <- rowSums(mydata.df))
  user  system elapsed 
 0.08    0.00    0.08
```

```
> system.time(mydata.df[mydata.df$X100 > 0 & mydata.df$X200 < 0,])
  user  system elapsed 
 0.72    0.05    0.76 
> system.time(mydata.df[which(mydata.df$X100 > 0 & mydata.df$X200 < 0),])
  user  system elapsed 
 0.25    0.00    0.25
```




(A5) Use of Hash Tables for Frequent Lookups 21/66 on Large Data

- The implementation of `lists` in R is not optimized for lookup; it incurs $O(N)$ time complexity to perform a lookup on a list of N elements.
- A hash table's (雜湊表) lookup incurs $O(1)$ time complexity. (R package `hash`)

```
> n <- 1E6
> mydata <- rnorm(n)
> mydata.ls <- as.list(mydata)
> names(mydata.ls) <- paste("X", c(1:n), sep="")
> head(mydata.ls, 3)
$X1
[1] -0.1038026

$X2
[1] -0.09553649

$X3
[1] -0.7474468

> library(hash)
> mydata.hs <- hash(names(mydata.ls), mydata)
> str(mydata.hs)
Formal class 'hash' [package "hash"] with 1 slot
..@ .xData:<environment: 0x00000000560a3fe8>
>
> #create lookups
> id <- sample(1:n, size=1000, replace=T)
> lookups <- paste("X", id, sep="")
```

雜湊表(hash table)是一種可以快速處理資料新增、搜尋及刪除的資料結構。利用資料的鍵值(key)直接對應至儲存位置的方法，雜湊表可以在幾次的資料比對後就完成資料加入、搜尋及刪除的動作。

```
> #comparison
> comptime.list <- sapply(lookups,
+   FUN = function(x){
+     system.time(mydata.ls[[x]])[3]}
+ )
> sum(comptime.list)
[1] 8.97
>
>
> comptime.hash <- sapply(lookups,
+   FUN = function(x){
+     system.time(mydata.hs[[x]])[3]}
+ )
> sum(comptime.hash)
[1] 0.17
```



(A6) Seeking Fast Alternative Packages in CRAN

22/66

- `fastcluster`: using an optimized C++ code that improves the speed significantly compared to the routines implemented in `hclust`.
- `fastmatch`: a faster version of base R's `match` function
- `RcppEigen`: a faster version of linear modeling `lm`
- `data.table`: This offers faster data manipulation operations compared to the standard `data.frame` operations
- `dplyr`: This offers a set of tools to manipulate data frame-like objects efficiently

```
> n <- 1E4
> p <- 100
> mydata <- matrix(rnorm(n*p), nrow=n, ncol=p)
> mydata.dist <- dist(mydata)
> system.time(hc.obj1 <- hclust(mydata.dist))
  user  system elapsed 
 4.46   0.20   4.66 
> 
> library(fastcluster)
> system.time(hc.obj2 <- hclust(mydata.dist))
  user  system elapsed 
 2.59   0.11   2.70
```

See also: <http://www.rdocumentation.org/>



(B) Using Compiled Code for Greater Speed

23/66

(B1) Compiling R code before execution.

- **compiler** package: allow to compile R code beforehand and save R a step or two when we execute the code.
- **jit** package: enable just-in-time (JIT) compilation of a block of R code.

(B2) Using compiled languages such as C and use them from within R.


- **inline** package : allows to embed C, C++, Objective-C, Objective-C++, and Fortran code within R.

(B3) Calling external compiled code

- R provides a few interfaces to call the external compiled code:
.C(), **.Fortran()**, **.Call()**, **.External()**
- **Rcpp** package : provides a convenient, higher-level API to the **.Call()** interface for C++ code.

(C) Using GPUs to Run R Even Faster 24/66

- **Graphics Processing Unit (GPU)**, known as a graphics card.
- To achieve real-time rendering, most GPU computations are done in a highly parallel manner, with many more cores than CPUs - a modern GPU might have more than 2,000 cores. Given that one core can run multiple threads, it is possible to run tens of thousands of parallel threads on a GPU.
- Will need an NVIDIA GPU with CUDA capabilities.
 - CUDA(Compute Unified Device Architecture)是NVIDIA 的平行運算架構，可運用繪圖處理單元(GPU) 的強大處理能力，大幅增加運算效能。



<http://www.nvidia.com.tw/object/what-is-gpu-computing-tw.html>



Example of Using GPUs to Run R ^{25/66}

- Calculate Kendall correlations on random datasets having 100 variables with a varying number of observations from 100, 200, ... to 500 records in order to observe the speedup in comparison to the CPU version.


```
library(gputools)
A <- lapply(c(1:5), function(x){
  matrix(rnorm((x*1E2) * 1E2), 1E2, (x*1E2))

  cpu.k.time <- sapply(A, function(x){
    system.time(cor(x, "kendall"))[[3]]})

  gpu.k.time <- sapply(A, function(x) {
    system.time(gpuCor(x, "kendall"))[[3]]})
```

```
> str(A)
List of 5
 $ : num [1:100, 1:100] 0.0318 1.3748 -0.0756 -0.1787 -0.4244 .
 $ : num [1:100, 1:200] -0.7577 -0.0139 0.8256 1.1097 0.1976 ..
 $ : num [1:100, 1:300] -0.838 0.398 0.934 0.13 1.071 ...
 $ : num [1:100, 1:400] 0.71 1.9 0.975 -0.341 1.686 ...
 $ : num [1:100, 1:500] -0.89 -1.378 -0.517 -0.552 1.69 ...
```

Tested on an NVIDIA GRID K520 GPU from AWS.





R Packages Related to GPUs

26/66

- [gputools](#): provides some data-mining algorithms which are implemented using a mixture of nVidia's CUDA language and cublas library.
- [cudaBayesreg](#): implements the rhier Linear Model from the [bayesm](#) package using nVidia's CUDA language and tools to provide high-performance statistical analysis of fMRI voxels.
- [rgpu](#): aims to speed up bioinformatics analysis by using the GPU.
- [gcbd](#): implements a benchmarking framework for BLAS and GPUs.
- [OpenCL](#): provides an interface from R to OpenCL permitting hardware- and vendor neutral interfaces to GPU programming.
- [HiPLARM](#): provide High-Performance Linear Algebra for R using multi-core and/or GPU support using the PLASMA / MAGMA libraries from UTK, CUDA, and accelerated BLAS.
- [gmatrix](#): enables the evaluation of matrix and vector operations using GPU coprocessors such that intermediate computations may be kept on the coprocessor and reused, with potentially significant performance enhancements by minimizing data movement.
- [gpuR](#): offers GPU-enabled functions for mirroring typical R syntax without the need to know OpenCL.



Factors that Affect the GPU's Performance

- GPUs work best for data parallel problems.
 - They are not suited for tasks that require large amounts of synchronization between threads.
- GPU's performance depends on the amount of data transferred between RAM and the GPU's memory.
 - The connection between the RAM and GPU's memory has a low bandwidth.
 - Good GPU programming should minimize this data transfer.
- 缺點:
 - Addressing these factors requires programming in the low-level GPU interfaces provided by RCUDA or OpenCL.
 - Packages 安裝稍為複雜。需硬體配合。
 - 目前較少分析分法: `gputools`: `gpuLm`, `gpuGlm`, `gpuCor`, `gpuDist`, `gpuHclust`, ...
 - 不一定比CPU快。



(II) Use Less RAM

(A) Simple Tweaks to Use Less RAM

(A1) Reusing objects without taking up more memory

- the **copy-on-modification** semantics of R's memory management. (`b <- a` # appears to make a copy of a and refer to it as b)

(A2) Removing intermediate data when it is no longer needed

(A3) Calculating values on the fly instead of storing them persistently

- On the fly computations produce intermediate data without creating variables that persist in the memory.
- Functions are a useful way to group related operations and automatically remove temporary variables when exiting the functions.

(A4) Swapping active and nonactive data

- save data to the disk to free up memory and reload them later when needed.



(B) Processing Large Datasets with Limited RAM

29/66

- (B1) Using memory-efficient data structures and smaller data types
- (B2) Sparse matrices
- (B3) Symmetric matrices
- (B4) Bit vectors
- (B5) Using memory-mapped files and processing data in chunks

■ R packages:

- **bigmemory**: **The Bigmemory Project** (Working with Massive Matrix-like Objects in R): **biganalytics**, **bigtabulate**, **synchronicity**, **bigalgebra**. (<http://www.bigmemory.org>)
- **biglm**: bounded memory linear and generalized linear models for data too large to fit in memory.
- **ff**:
 - supports more data types than **bigmemory**.
 - provides a more data frame-like memory-mapped format is required while dealing with heterogeneous data types.
 - memory-efficient storage of large data on disk and fast access functions.



(B1) Using Memory-efficient Data Structures and Smaller Data Types

30/66

Memory-efficient Data Structures

```
> print(object.size(logical(1e6)), units="MB")
3.8 Mb
> print(object.size(integer(1e6)), units="MB")
3.8 Mb
> print(object.size(numeric(1e6)), units="MB")
7.6 Mb
> print(object.size(complex(1e6)), units="MB")
15.3 Mb
> print(object.size(rep.int(NA_character_, 1e6)), units="MB")
7.6 Mb
> print(object.size(raw(1e6)), units="MB")
1 Mb
> print(object.size(vector("list", 1e6)), units="MB")
7.6 Mb
```

Smaller data types

```
> object.size(as.numeric(seq_len(1e6)))
8000040 bytes
> object.size(as.integer(seq_len(1e6)))
4000040 bytes
> strings <- rep.int(formatC(seq_len(1e4), width = 1000), 100)
> object.size(strings)
18480040 bytes
> factors <- factor(strings)
> object.size(factors)
14560400 bytes
```



(B2) Sparse Matrices

```
> library(Matrix)
> n <- rnorm(1e6)
> n[sample.int(1e6, 7e5)] <- 0
> m.dense <- Matrix(n, 1e3, 1e3, sparse = FALSE)
> m.sparse <- Matrix(n, 1e3, 1e3, sparse = TRUE)
> object.size(n)
8000040 bytes
> object.size(m.dense)
8001112 bytes
> object.size(m.sparse)
3605424 bytes
```

Sparse matrices are also very useful for binary data (TRUE/FALSE, 0/1, "yes"/"no", "hot"/"cold", and so on).

```
> logicalData <- sample(c(FALSE, TRUE), 1e6, TRUE, c(0.7, 0.3))
> m2.dense <- Matrix(logicalData, 1e3, 1e3, sparse = FALSE)
> m2.sparse <- Matrix(logicalData, 1e3, 1e3, sparse = TRUE)
> object.size(logicalData)
4000040 bytes
> object.size(m2.dense)
4001112 bytes
> object.size(m2.sparse)
2410576 bytes
```

Sparse data: it contains a lot of zeroes or empty values.

```
> m.dense[1:3,1:3]
3 x 3 Matrix of class "dgeMatrix"
      [,1] [,2] [,3]
[1,] 0.0000000 0 0.0000000
[2,] -0.7627943 0 0.0000000
[3,] 0.0000000 0 0.2869535

> m.sparse[1:3,1:3]
3 x 3 sparse Matrix of class "dgCMatrix"

[1,] . . .
[2,] -0.7627943 . .
[3,] . . 0.2869535
```

The sparse logical matrix is even more compact than the sparse numeric matrix,



(B3) Symmetric Matrices

- The **Matrix** package provides the **dspMatrix** class to efficiently store symmetric matrices and other efficient matrix-type data structures including triangular matrices and diagonal matrices.
 - The package makes it such that basic matrix operations, such as matrix multiplication (`%*%`), are applicable for both dense and sparse matrices.

```
> library(Matrix)
> data <- matrix(rnorm(1E5), 1E2, 1E3)
> A <- cor(data)
> isSymmetric(A)
[1] TRUE
> B <- as(A, "dspMatrix")
> object.size(A)
8000200 bytes
> object.size(B)
4005320 bytes
```


(B4) Bit Vectors

- Unlike logical values (Binary data) in R that take up four bytes (32 bits), **bit vectors** store each logical value using only one bit.
 - This reduces the memory consumption of logical values by a factor of 32.
 - Bit vectors cannot store the NA value, so they are not suitable for data that contains the NA values.
- **bit** package:
 - When dealing with large amounts of logical or binary data, bit vectors not only save memory but also provide a speed boost when they are operated on.

```
> # compare the sizes of a logical vector and the equivalent bit vector
> library(bit)
> a <- sample(c(TRUE, FALSE), 1e6, TRUE)
> object.size(a)
4000040 bytes
> b <- as.bit(a)
> object.size(b)
126344 bytes
```

(B5) Using Memory-mapped Files and Processing Data in Chunks 4/66

- Some datasets are still too large to fit in or be processed in the memory. One way to work with such large data is to store them on a disk in the form of **memory-mapped files** and load the data into the memory for processing one small chunk at a time.
- Some algorithms can easily be converted to compute on chunks of data, while others might require substantial effort to do so.

```
> setwd("D:/test-R")
> library(bigmemory)
> #object bm.obj stores a pointer to the new memory-mapped file bm.
> bm.obj <- big.matrix(1e9, 3, backingfile = "bm", backingpath = getwd())
> bm.obj
```

```
An object of class "big.matrix"
Slot "address":
<pointer: 0x00000000fd201b0>
```

```
# the new file bm has a size of 22 GB
#"bm.desc" was created to describe the data file.
# This is used to retrieve the memory-mapped file at
a later time such as
my.bm <- attach.big.matrix("bm.desc")
```

Example: A Chunked Computation

```
#fills in bm.obj with random numbers in 100 chunks of 10 million rows at a time
chunksize <- 1e7
start <- 1
while (start <= nrow(bm.obj)) {
  end <- min(start + chunksize - 1, nrow(bm.obj))
  chunksize <- end - start + 1
  bm.obj[start:end, 1] <- rpois(chunksize, 1e3)
  bm.obj[start:end, 2] <- sample(0:1, chunksize, TRUE, c(0.7, 0.3))
  bm.obj[start:end, 3] <- runif(chunksize, 0, 1e5)
  start <- start + chunksize
}
```

```
> dim(bm.obj)
[1] 1e+09 3e+00

> head(bm.obj)
      [,1] [,2]      [,3]
[1,] 1007   1 31038.563
[2,]  970   0  4738.024
[3,] 1027   0 65442.558
[4,] 1012   0 67877.496
[5,] 1027   0 71838.828
[6,] 1032   0 29757.398
```

When the subsetting operator `[,]` is used, **bigmemory** automatically loads the relevant portions of the data into the RAM and removes portions that are no longer needed.



A Chunked Computation to Find the Standard Deviation of Each Column

36/66

```
> col.sums <- numeric(3)
> chunksize <- 1e7
> start <- 1
> while (start <= nrow(bm.obj)) {
+   end <- min(start + chunksize - 1, nrow(bm.obj))
+   col.sums <- col.sums + colSums(bm.obj[start:end, ])
+   start <- start + chunksize
+ }
> col.means <- col.sums/nrow(bm.obj)
> col.means
[1] 1000.001289      0.299985 49999.209345
```

Calling `sd(bm[1,])` might not work, as even a single column of data can exceed available memory.

```
> col.sq.dev <- numeric(3)
> start <- 1
> while (start <= nrow(bm.obj)) {
+   end <- min(start + chunksize - 1, nrow(bm.obj))
+   sq.dev <- sq.dev + rowSums((t(bm.obj[start:end, ])- col.means)^2)
+   start <- start + chunksize
+ }
> col.var <- sq.dev/(nrow(bm.obj) - 1)
> col.sd <- sqrt(col.var)
> col.sd
[1] 31.623471      0.458251 28867.409206
```

```
> library(biganalytics)
> colsd(bm.obj)
[1] 31.623471      0.458251 28867.409206
```



Example: The Bigmemory Project

- 例子: Airline on-time performance data from the 2009 JSM Data Expo
- 資料大小: airline.csv: 約 **11 GB** (120 million rows and 29 columns).
- 運算平台: A laptop with only 4 GB of RAM
- 資料讀取結果:

```
> library(bigmemory)
> library(biganalytics)
> x <- read.big.matrix("airline.csv", type="integer", header=TRUE,
  ...) #約 25分鐘。
> summary(x) #約 3-4 分鐘
```

- 統計分析結果:

`lm()` 需要超過10GB的記憶體

```
> blm <- biglm.big.matrix(ArrDelay ~ Age + Year, data=x)
# 3分鐘 + 幾百MB記憶體
> summary(blm)
Large data regression model: biglm(formula = formula, data = data,
  ...)
Sample size = 84216580
...
```

- 結論: 不是每個統計方法in R都有相對應的big版本。

<http://www.bigmemory.org>

平行化計算？





Amdahl's Law (阿姆達爾定律)

- Most algorithms fall somewhere in between with some steps that must run **serially** and some that can run in **parallel**.
- Amdahl's law provides a way to estimate the best attainable performance gain when you convert a code from serial to parallel execution:

$$T(n) = T(1)(P + (1-P)/n)$$

- $T(n)$ is the time taken to execute the task using n parallel processes
- P is the proportion of the whole task that is strictly serial
- The theoretical best possible speed up of the parallel algorithm is:

$$S(n) = T(1) / T(n) = 1 / (P + (1-P)/n)$$


Example: given a task that takes 10 seconds to execute on one processor, where half of the task can be run in parallel, then the best possible time to run it on four processors is

$$T(4) = 10(0.5 + (1-0.5)/4) = 6.25 \text{ seconds.}$$

The theoretical best possible speed up of the parallel algorithm with four processors is

$$1 / (0.5 + (1-0.5)/4) = 1.6x.$$

How Much Can Be Parallelized?^{40/66}




<https://en.wikipedia.org/wiki/File:AmdahlsLaw.svg>

Data Parallelism vs Task Parallelism


41/66

- In data parallelism, a dataset is divided into multiple partitions.
- Different partitions are distributed to multiple processors, and the **same task** is executed on each partition of data.

Serial execution



Data parallel execution




Source: Aloysius Lim, and William Tjhi, R High Performance Programming, Packt Publishing, January 30, 2015.



Data Parallelism vs Task Parallelism

42/66

- In task parallelism, tasks are distributed to and executed on different processors in parallel.
 - The tasks on each processor might be the same or different, and the data that they act on might also be the same or different.
 - The key difference from data parallelism: the data is not divided into partitions.
- Example (data parallelism): A random forest is a collection of decision trees built independently on the same data. During the training process for a particular tree, a random subset of the data is chosen as the training set, and the variables to consider at each branch of the tree are also selected randomly. Hence, even though the same data is used, the trees are different from one another.
- Example (task parallelism): Training of a random forest model. In order to train a random forest of say 100 decision trees, the workload could be distributed to a computing cluster with 100 processors, with each processor building one tree. All the processors perform the same task on the same data (or exact copies of the data), but the data is not partitioned.



See Parallel R by Q. Ethan McCallum and Stephen Weston.



Get Parallel Results Without Explicitly Doing Parallel Programming

43/66

- **Plyr** package
 - **aaply**: like **apply**, but with an option to parallelize.
- **Foreach** package
 - Lets you write for loops that can be parallelized (assuming only effect of code is the return value)
- **Multicore** package
 - If you have a machine with a multi-core processor, the multicore apply functions (e.g. **mclapply**) can farm some set of the data on which the **apply** is being called to the other core to speed up the process.
 - Multicore also has functions for more detailed parallel programming, which can make things fast but forces you to confront **concurrency** (並發性來源) issues.



(IV) Offloading Data Processing to Database Systems

44/66

- Extracting data into R versus processing data in a database
 - In some situations, this might not be efficient to manipulate data in R by moving all the data into R whether in memory or on a disk, on a single computer or on a cluster.
 - We can process the data in the database and retrieve only the results into R, which are usually much smaller than the raw data.
 - At the foundation of all in-database tools is the SQL language (relational databases). SQL is a powerful and flexible language used to manipulate data in a database.
- Use an R package as an API to the data warehouse
 - R packages: **RJDBC**, **RODBC**, **RPostgreSQL**, **RMySQL**, and **ROracle**.
- Converting R expressions to SQL:
 - R packages: **dplyr**, **PivotalR**, **MonetDB.R** and **SciDB**.

Example: Running Statistical and Machine Learning Algorithms in a Database

45/66

- Open source project, **MADlib** (<http://madlib.net/>)
 - brings advanced statistics and machine learning capabilities to PostgreSQL databases, including
 - descriptive statistics, hypothesis tests, array arithmetic, probability functions, dimensionality reduction, linear models, clustering models, association rules, and text analysis.

```
db.drv <- PostgreSQL()  
db.conn <- dbConnect(db.drv, host = "hostname",  
                    port = 5432, dbname = "rdb",  
                    user = "ruser",  
                    password = "rpassword")
```

```
dbGetQuery(db.conn,  
  "SELECT *  
  FROM madlib.assoc_rules(  
    0.001, -- support  
    0.01, -- confidence  
    'trans_id', -- tid_col  
    'item_id', -- item_col  
    'trans_items', -- input_table  
    'public', -- output_schema  
    TRUE -- verbose  
  );")
```

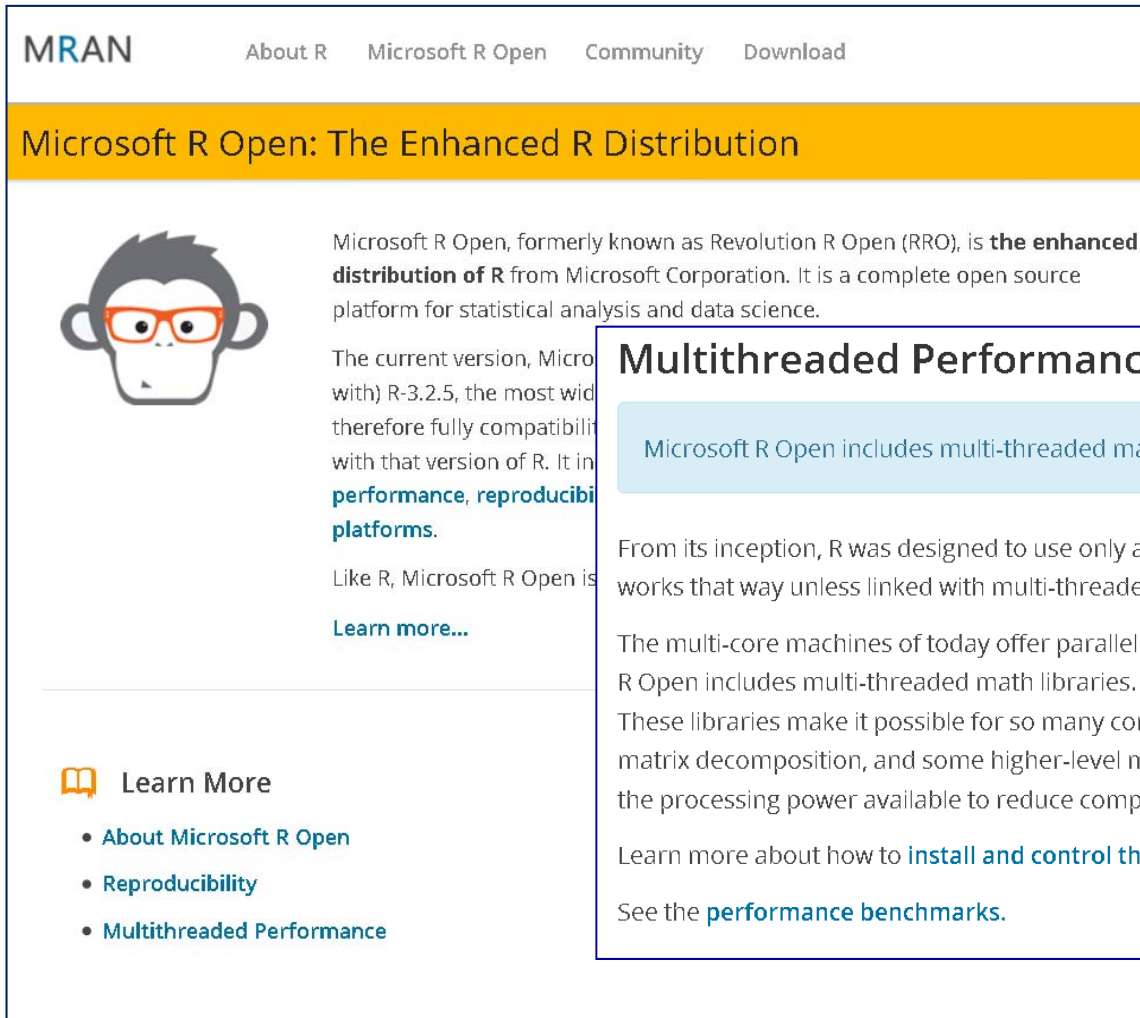


The screenshot shows the MADlib website homepage. At the top, there is a navigation bar with links for Home, Product, Documentation, and Community, along with a red Download button. The main header features the MADlib logo and the text "Apache MADlib (incubating): Big Data Machine Learning in SQL for Data Scientists". Below this, there are three columns of text: "Open source, commercially friendly Apache license", "Supports PostgreSQL, Greenplum Database™, and Apache HAWQ (incubating)", and "Powerful analytics for big data". A "Read More" link is positioned to the right of the third column. At the bottom of the page, it says "MADlib 1.9 Release (GA)".

Linux (Red Hat), Mac OS X


(VI) Other Consideration

(A) Microsoft R Open



MRAN About R Microsoft R Open Community Download

Microsoft R Open: The Enhanced R Distribution



Microsoft R Open, formerly known as Revolution R Open (RRO), is **the enhanced distribution of R** from Microsoft Corporation. It is a complete open source platform for statistical analysis and data science.

The current version, Microsoft R Open (with R-3.2.5, the most widely used version), is therefore fully compatible with that version of R. It includes **performance, reproducibility, and cross-platform support**.

Like R, Microsoft R Open is open source.

[Learn more...](#)

Learn More

- [About Microsoft R Open](#)
- [Reproducibility](#)
- [Multithreaded Performance](#)

Microsoft R Application Network

The Microsoft R Portal

Multithreaded Performance

Microsoft R Open includes multi-threaded math libraries to improve the performance of R.

From its inception, R was designed to use only a single thread (processor) at a time. Even today, R works that way unless linked with multi-threaded BLAS/LAPACK libraries.

The multi-core machines of today offer parallel processing power. To take advantage of this, Microsoft R Open includes multi-threaded math libraries.

These libraries make it possible for so many common R operations, such as matrix multiply/inverse, matrix decomposition, and some higher-level matrix operations, to compute in parallel and use all of the processing power available to reduce computation times.

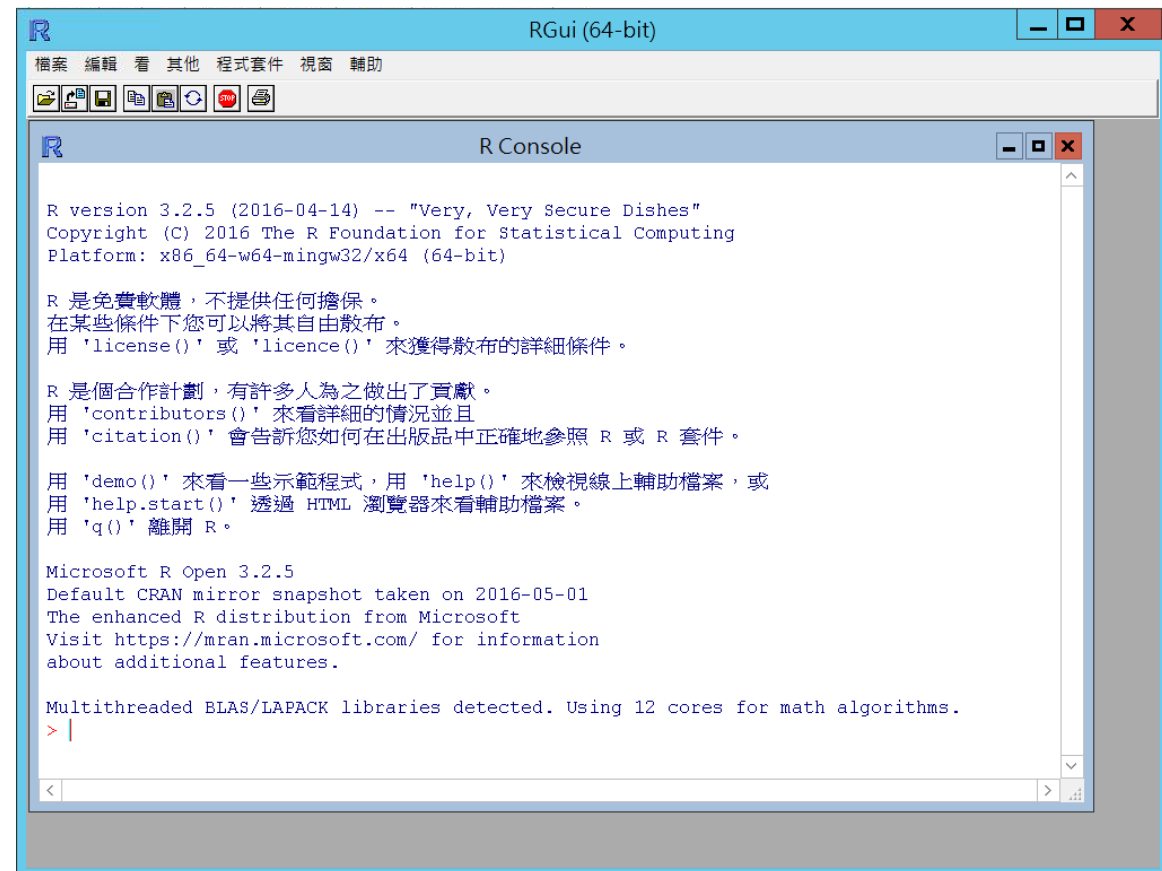
Learn more about how to [install and control the number of threads](#).

See the [performance benchmarks](#).

<https://mran.revolutionanalytics.com/documents/rro/multithread/>

Multithreading on Windows Math Kernel Library (MKL)

47/66



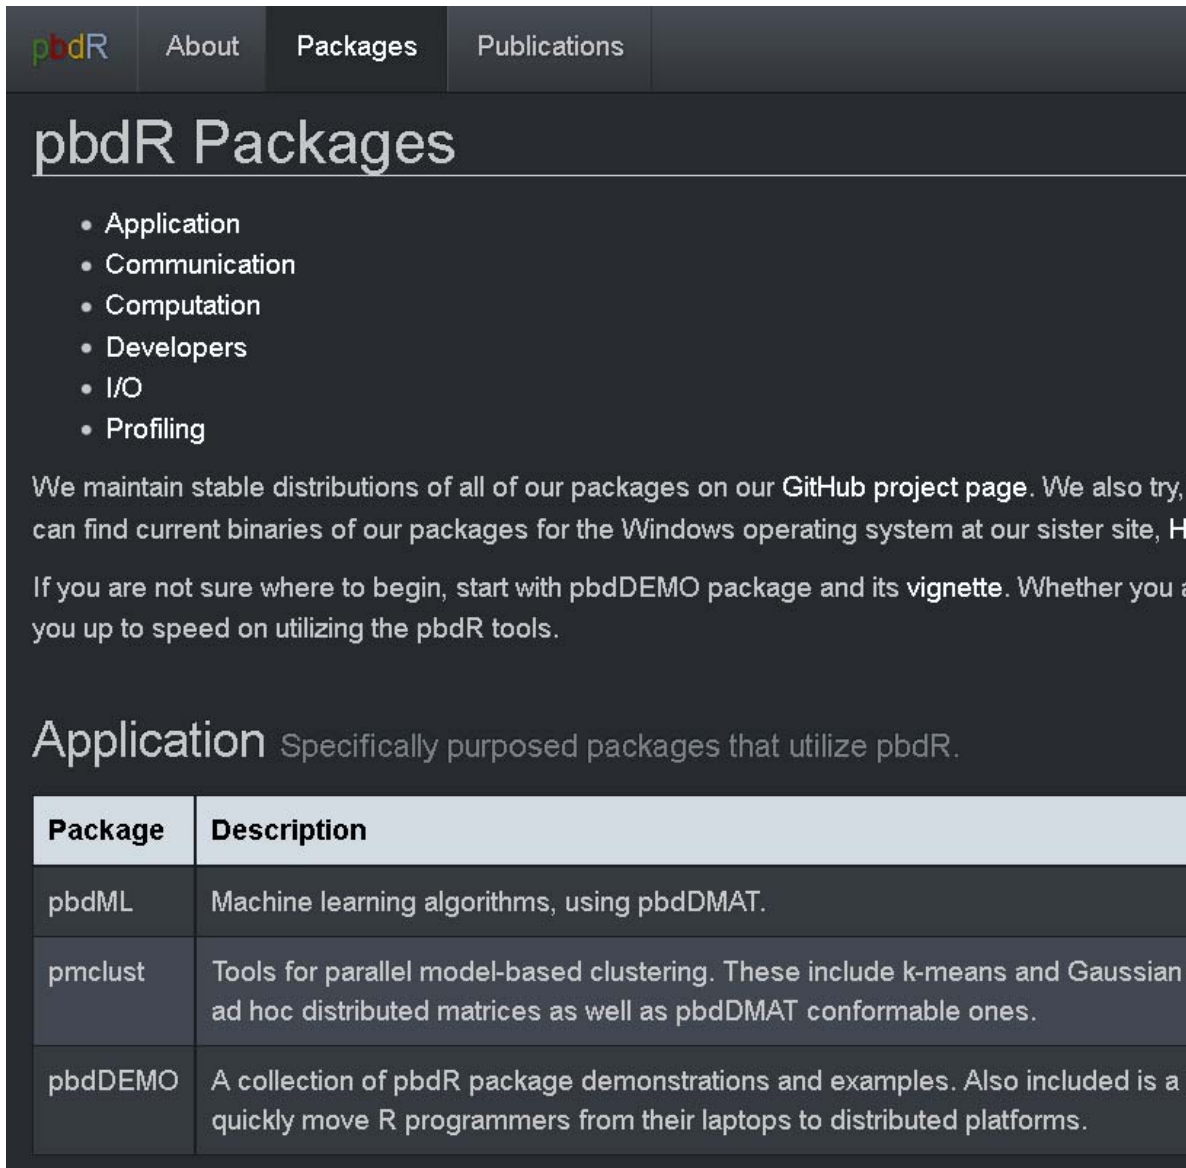
Elapsed time (**Relative Performance**)

	R-3.2.5 (1 thread)		Microsoft R Open-3.2.5 (1 thread)		Microsoft R Open-3.2.5 (4 threads)		Microsoft R Open-3.2.5 (8 threads)	
Matrix multiplication	140.63	1	14.11	9.97	3.27	43.01	1.89	74.41
Cholesky factorization	20.27	1	0.42	48.26	0.35	57.91	0.36	56.31
QR decomposition	4.36	1	2.78	1.57	2.59	1.68	2.63	1.66
Singular value decomposition	52.07	1	1.91	27.26	1.89	27.55	1.90	27.41
Principal component analysis	40.46	1	3.25	12.45	2.67	15.15	2.55	15.87
Linear discriminant analysis	10.14	1	2.36	4.30	2.04	4.97	2.03	5.00

Table Modified from: <https://mran.revolutionanalytics.com/documents/rro/multithread/>

(B) Programming with Big Data in R

49/66



pbdR About Packages Publications

pbdR Packages

- Application
- Communication
- Computation
- Developers
- I/O
- Profiling

We maintain stable distributions of all of our packages on our [GitHub project page](#). We also try, where possible, to find current binaries of our packages for the Windows operating system at our sister site, [HPC](#).


If you are not sure where to begin, start with pbdDEMO package and its [vignette](#). Whether you are a beginner or an expert, you are up to speed on utilizing the pbdR tools.

Application

Specifically purposed packages that utilize pbdR.

Package	Description
pbdML	Machine learning algorithms, using pbdDMAT.
pmclust	Tools for parallel model-based clustering. These include k-means and Gaussian mixture models on large and ad hoc distributed matrices as well as pbdDMAT conformable ones.
pbdDEMO	A collection of pbdR package demonstrations and examples. Also included is a vignette to help you quickly move R programmers from their laptops to distributed platforms.

<http://r-pbd.org/>



pbdR Relationships to Libraries


a set of R packages for large
distributed computing and profiling.
programming
designed to help use R for
analyzing really big data on high-
performance computing clusters.



(VI) R and Big Data

- How do analysts use big data:
 - class 1: extract data:
 - extract subset, sample, or summary from a big data source.
 - note that the subset might itself be quite large.
 - class 2: compute on the parts:
 - repeat computations for many subgroups of the data and combine the results.
 - (analogous to [map-reduce/split-apply-combine](#))
Can use computing clusters to parallelize analysis.
 - class 3: compute on the whole:
 - the problems of using all of the data at once are irretrievably big.
 - they must be run at scale within the data warehouse.

Hadoop Techniques



- Setup Hadoop Environment (or Using Commercial Services)
 - **Linux** OS, (or VMware to host Ubuntu within the Windows OS.)
 - Understanding Hadoop features: the **Hadoop Distributed File System (HDFS)** and **MapReduce** architecture.
 - Writing Hadoop MapReduce Programs using **Java**.
- Integrating R and Hadoop:
 - Learning **RHadoop**: performing data analytics with the Hadoop platform via R functions.
 - Learning **RHIPE**: R and Hadoop Integrated Programming Environment (RHIPE) is widely used for performing Big Data analysis with Divide and Recombine (D&R) analysis.
 - Learning **Hadoop streaming**: Hadoop streaming allows you to create and run MapReduce jobs with any executable or R script as the Mapper and/or Reducer. See R package **HadoopStreaming**.
- R packages for importing and exporting data from various DBs: **RMySQL**, **Xlsx**, **RMongo**, **RSQLite**, **RPostgreSQL**, **RHive**, **RHBase**, **RHDFS**

Amazon Web Services

52/66

關於 AWS

AWS 免費方案

常見問答集

相關連結

什麼是雲端運算?

AWS 入門

AWS 產品與服務

AWS 上執行網站

AWS 新帳戶可以獲得 AWS 免費方案長達 12 個月的使用權

現在開始

AWS 免費方案

Amazon Web Services (AWS) 免費方案旨在提供您 AWS 雲端服務的實作經驗。AWS 免費方案包含 AWS 註冊日期後 12 個月提供免費方案的服務，以及不會在 AWS 免費方案期限 12 月結束後自動到期的其他服務。

建立 AWS 帳戶後，您可以在特定用量限制內，免費使用下列任何產品和服務。

您可以立即開始，並依照下列步驟自動享有 AWS 免費方案的各項優勢：

1. 註冊 AWS 帳戶。
2. 輸入您的帳單地址和信用卡資訊。除非您的用量超過免費方案限制，否則無須付費。
3. 選取下列任一產品，開始使用 AWS 雲端服務。

排序方式: 特色產品

Amazon EC2

雲端中可調規模的運算容量。
[進一步了解](#)

750 小時 / 月的 Linux、RHEL 或 SLES t2.micro
750 小時 / 月的 Windows t2.micro 執行個體
例如，一個月執行 1 個執行個體，或者半個月執行 2 個執行個體。
註冊後 12 個月到期。

Amazon S3

高度可擴展性、可靠、低延遲的資料儲存體基礎設施。
[進一步了解](#)

5 GB 的標準儲存
20,000 個 Get 請求
2,000 個 Put 請求
註冊後 12 個月到期。

開始免費使用 AWS - 按使用量付費

建立免費帳戶

Region Unsupported

Machine Learning is not available in US West (Oregon).

Supported Regions

EU (Ireland)
US East (N. Virginia)

- **m1.xlarge**: AWS servers: 4 CPUs and 15 GB of RAM each, with cluster sizes ranging from 4 to 32 core nodes.
- **cr1.8xlarge**: the largest memory size that an Amazon EC2 node can have is 244 gigabytes.

- Amazon **EC2**: Amazon Elastic Compute Cloud
- Elastic MapReduce (**EMR**) allows to rent and run a Hadoop cluster on an hourly basis.
- An Amazon Machine Image (**AMI**) defines which operating system is installed on your EC2 instance and which software is included.

Amazon Web Services

53/66

The screenshot shows the AWS EC2 Management Console interface. The left sidebar contains navigation links for EC2 Dashboard, Events, Tags, Reports, Limits, INSTANCES, IMAGES, ELASTIC BLOCK STORE, and NETWORK & SECURITY. The main content area displays a table of EC2 instances. A text box with a dashed red border highlights the following text:

With AWS, you can get a server with up to 244 GB of main memory and up to 40 CPUs; no more limitations by hardware and computation time for your R-based analyses. The Amazon Linux AMI is a good starting point for setting up your own analysis environment.

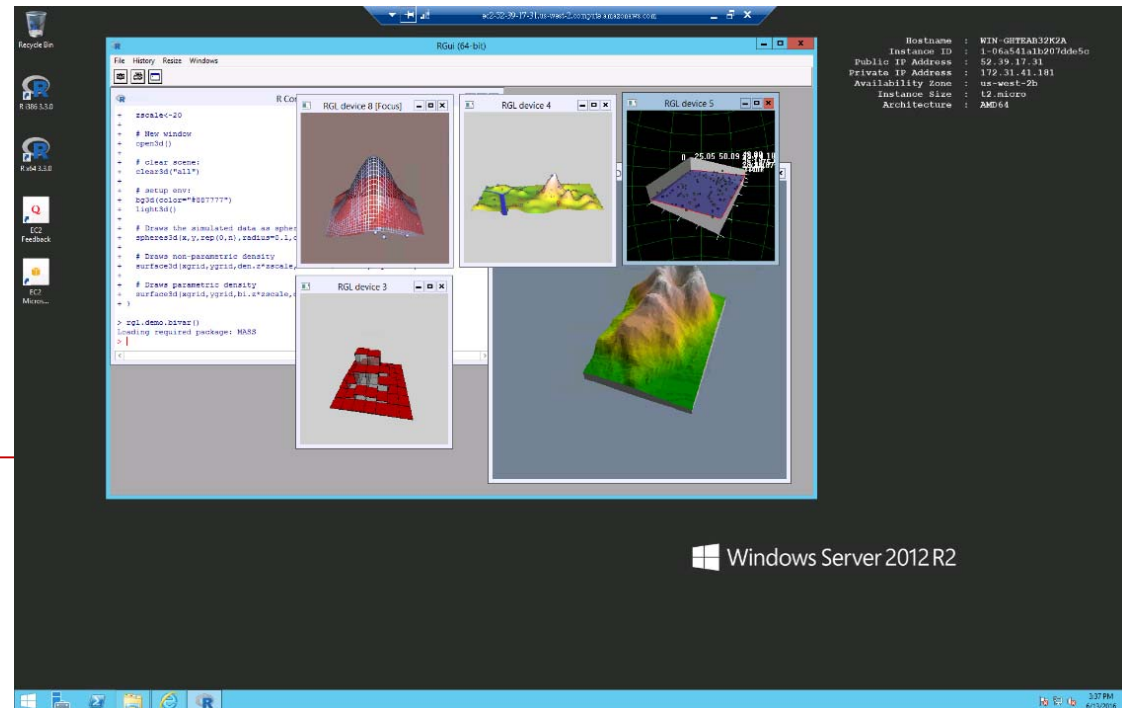
Name	Instance ID	Instance Type	Availability Zone	Instance State	Status Checks	Alarm Status
	i-06a541a1b207dde5c	t2.micro	us-west-2b	running	2/2 checks ...	None
	i-07162a3bc3509e9f6	t2.micro	us-west-2a	terminated		None
	i-0fc3e9034fbac0372	t2.micro	us-west-2b	terminated		None

At the bottom of the console, the selected instance details are shown: Instance: i-06a541a1b207dde5c, Public DNS: ec2-52-39-17-31.us-west-2.compute.amazonaws.com.

Setting Up an AWS Instance for R ^{54/66}

- Uploading data to HDFS
- Analyzing HDFS data with RHadoop
- R code for MapReduce

```
mapper <- function(keys, values) {  
  ...  
}  
  
reducer <- function(key, values) {  
  ...  
}  
  
job <- mapreduce(input = "/data/myinput",  
                 input.format = input.format,  
                 output = "/data/myoutput",  
                 map = mapper, reduce = reducer)  
  
results <- from.dfs(job)
```



How to set up an AWS instance for R

- <https://blogs.aws.amazon.com/bigdata/post/Tx3IJSB6BMHWZE5/Running-R-on-AWS>
- <http://www.r-bloggers.com/setting-up-an-aws-instance-for-rstudio-opencpu-or-shiny-server/>
- https://www.r-project.org/conferences/useR-2010/tutorials/Zolot_tut.pdf

Microsoft Azure



Microsoft Azure：雲端計算...

Microsoft Corporation (US) | <https://azure.microsoft.com/zh-tw>

Microsoft Azure

銷售專線 0800-00-88-33#2#9 | 我的帳戶 | 入口網站 | 搜尋

為何選擇 Azure? | 產品 | 文件 | 定價 | 合作夥伴 | 部落格 | 資源 | 支援

免費帳戶 >

現代化商務的雲端

加快腳步
節省成本
在內部部署環境整合應用程式和資料

免費試用 >


在研發/測試及生產環境中執行您的 SAP HANA 應用程式 >

開始使用

觀看解說如何快速開始使用 Microsoft Azure 的三分鐘影片。


探索更多 >

熱門產品


<p>Virtual Machines</p>  <p>只要數分鐘的時間，就能完成 Windows 與 Linux 虛擬機...</p>	<p>App Service</p>  <p>為任何平台及任何裝置建立 Web 和行動應用程式</p>	<p>SQL Database</p>  <p>以服務方式提供受管理的關聯式 SQL Database</p>	<p>Storage</p>  <p>耐用、具高可用性並可大幅調整的雲端儲存體</p>
<p>Cloud Services</p>  <p>建立高可用性、可無限擴充的雲端應用程式及 API</p>	<p>Document DB</p>  <p>受管理的 NoSQL 文件資料庫即服務</p>	<p>Active Directory</p>  <p>同步處理內部部署目錄和啟用單一登入</p>	<p>HDInsight</p>  <p>佈建雲端 Hadoop、Spark、R Server、HBase 及 Storm 叢...</p>

The Four Main Stages of Hadoop MapReduce Data Processing

56/66



The HDFS and MapReduce architecture



- The loading of data into HDFS
- The execution of the Map phase
- Shuffling and sorting
- The execution of the Reduce phase

- Map: 從主節點(master node)輸入一組input，此input是一組key/value，將這組輸入切成好幾個小的子部分，分散到各個工作節點(worker nodes)去做運算
- Reduce: 主節點(master node)收回處理完的子部分，將子部分重新組合產生輸出

Example: The Steps to Run a MapReduce Job with Hadoop

57/66

Map.java: the Map class for the word count Mapper.

```
// Defining package of the class
package com.PACKT.chapter1;

// Importing java libraries
import java.io.*;
import java.util.*;
import org.apache.hadoop.io.*;
import org.apache.hadoop.mapred.*;

// Defining the Map class
public class Map extends MapReduceBase
    implements Mapper<LongWritable,
        Text,
        Text,
        IntWritable>{

//Defining the map method - for processing
// problem specific logic
public void map(LongWritable key,
    Text value,
    OutputCollector<Text>
        ...

// Emitting the (key,value) pair with
// problem specific logic
output.collect(new Text(st.nextToken()),
    new IntWritable(1));
}
```

Reduce.java: the Reduce class for the word count Reducer.

```
// Defining package of the class
package com.PACKT.chapter1;

// Importing java libraries
import java.io.*;
import java.util.*;
import org.apache.hadoop.io.*;
import org.apache.hadoop.mapred.*;

// Defining the Reduce class
public class Reduce extends MapReduceBase
    implements Reducer<Text,
        IntWritable,
        Text,
        IntWritable>{

// Defining the reduce method
// generated output of Map phase
public void reduce(Text key,
    Iterator<IntWritable> values,
    OutputCollector<Text>
        ...

// Emitting the (key,value) pair with
// problem specific logic
output.collect(key, new IntWritable(1));
}
```

WordCount.java: the task of Driver in the Hadoop MapReduce Driver main file.

```
//Defining package of the class
package com.PACKT.chapter1;

// Importing java libraries
import java.io.*;
import org.apache.hadoop.fs.*;
import org.apache.hadoop.io.*;
import org.apache.hadoop.mapred.*;
import org.apache.hadoop.util.*;
import org.apache.hadoop.conf.*;

//Defining wordcount class for job configuration
// information
public class WordCount extends Configured implements Tool{

    ...

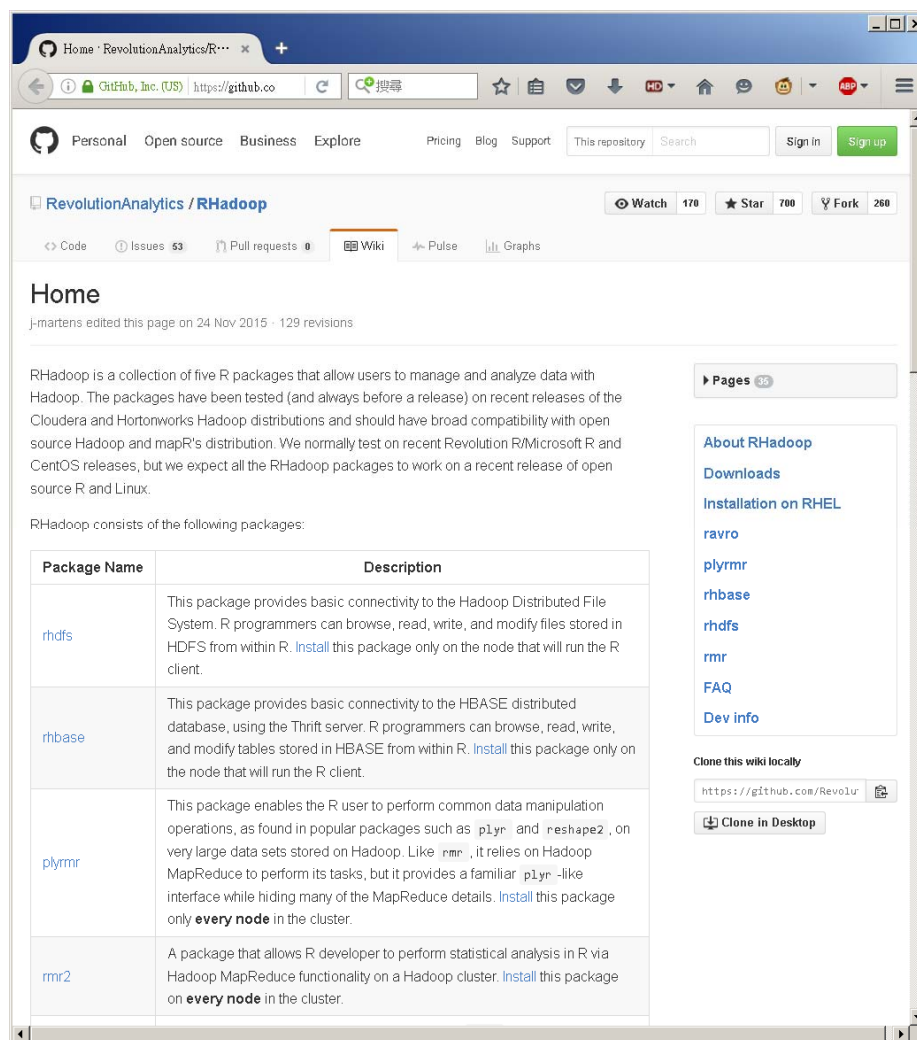
// For submitting the configuration object
JobClient.runJob(conf);

return 0;
}

// Defining the main() method to start the execution of //
// the MapReduce program
public static void main(String[] args) throws Exception {
    int exitCode = ToolRunner.run(new WordCount(), args);
    System.exit(exitCode); } }
```

RHadoop

58/66



Home · RevolutionAnalytics/RHadoop

Watch 170 Star 700 Fork 260

Code Issues 53 Pull requests 0 Wiki Pulse Graphs

Home

j-martens edited this page on 24 Nov 2015 · 129 revisions

RHadoop is a collection of five R packages that allow users to manage and analyze data with Hadoop. The packages have been tested (and always before a release) on recent releases of the Cloudera and Hortonworks Hadoop distributions and should have broad compatibility with open source Hadoop and mapR's distribution. We normally test on recent Revolution R/Microsoft R and CentOS releases, but we expect all the RHadoop packages to work on a recent release of open source R and Linux.

RHadoop consists of the following packages:

Package Name	Description
rhdfs	This package provides basic connectivity to the Hadoop Distributed File System. R programmers can browse, read, write, and modify files stored in HDFS from within R. Install this package only on the node that will run the R client.
rhbase	This package provides basic connectivity to the HBASE distributed database, using the Thrift server. R programmers can browse, read, write, and modify tables stored in HBASE from within R. Install this package only on the node that will run the R client.
plymr	This package enables the R user to perform common data manipulation operations, as found in popular packages such as <code>plyr</code> and <code>reshape2</code> , on very large data sets stored on Hadoop. Like <code>rmm</code> , it relies on Hadoop MapReduce to perform its tasks, but it provides a familiar <code>plyr</code> -like interface while hiding many of the MapReduce details. Install this package only every node in the cluster.
rmm2	A package that allows R developer to perform statistical analysis in R via Hadoop MapReduce functionality on a Hadoop cluster. Install this package on every node in the cluster.

Pages 35

- About RHadoop
- Downloads
- Installation on RHEL
- [ravro](#)
- [plymr](#)
- [rhbase](#)
- [rhdfs](#)
- [rmm](#)
- [FAQ](#)
- [Dev info](#)

Clone this wiki locally

<https://github.com/RevolutionAnalytics/RHadoop>

[Clone in Desktop](#)

See Also [mapReduce](#) package.

```
map <- function(k,lines) {
  words.list <- strsplit(lines, '\\s')
  words <- unlist(words.list)
  return( keyval(words, 1) )
}

reduce <- function(word, counts) {
  keyval(word, sum(counts))
}

wordcount <- function (input, output=NULL) {
  mapreduce(input=input, output=output,
    input.format="text", map=map,
    reduce=reduce)
}

## read text files from folder wordcount/data
## save result in folder wordcount/out
## Submit job

hdfs.root <- 'wordcount'
hdfs.data <- file.path(hdfs.root, 'data')
hdfs.out <- file.path(hdfs.root, 'out')
out <- wordcount(hdfs.data, hdfs.out)

## Fetch results from HDFS
results <- from.dfs(out)
results.df <- as.data.frame(results,
  stringsAsFactors=F)
colnames(results.df) <- c('word', 'count')
head(results.df)
```



MapReduce in R

59/66

Logistic Regression in R

```
model <- glm(response ~., family=binomial(link='logit'), data=mydata)
```

Using Mapreduce Framework in R

```
lr.map <- function(., M){  
  Y <- M[,1]  
  X <- M[,-1]  
  keyval(1, Y * X * g(-Y *  
    as.numeric(X %*% t(plane))))  
}  
  
lr.reduce <- function(k, Z){  
  keyval(k, t(as.matrix(  
    apply(Z,2,sum))))  
}
```

```
logistic.regression <- function(input,  
  iterations, dims, alpha){  
  
  plane = t(rep(0, dims))  
  g <- function(z) 1/(1 + exp(-z))  
  for(i in 1:iterations){  
    gradient <- values(from.dfs(  
      mapreduce(input,  
        map = lr.map,  
        reduce = lr.reduce,  
        combine = T)))  
    plane = plane + alpha * gradient  
  }  
  plane  
}
```

logistic regression by gradient descent

Source: <https://github.com/RevolutionAnalytics/rmr2/blob/master/docs/tutorial.md>



References

- Large Datasets and You:
 - <http://www.mattblackwell.org/files/papers/bigdata.pdf>
- FasterR! HigherR! StrongerR! - A Guide to Speeding Up R Code for Busy People:
 - <http://www.noamross.net/blog/2013/4/25/faster-talk.html>
- Scalable Strategies for Computing with Massive Data:
 - <http://www.slideshare.net/fullscreen/joshpaulson/big-memory/1>
- Taking R to the Limit, Part II: Working with Large Datasets:
 - http://www.bytemining.com/wp-content/uploads/2010/08/r_hpc_II.pdf
 - <https://github.com/rstudio/webinars/blob/master/14-Work-with-big-data/14-Work-with-big-data.pdf>
- 20140317 MLDM Monday - 使用 RHadoop 做巨量資料分析
 - <https://www.youtube.com/watch?v=vmIHqe8JSXg>
- MapReduce and R: Short example on regression / forecasting
 - <https://www.youtube.com/watch?v=Q8kmAfpwAJQ>
- RDataMining.com: R and Data Mining, Building an R Hadoop System
 - <http://www.rdatamining.com/big-data/rhadoop>
- MapReduce in R
 - <https://github.com/RevolutionAnalytics/rmr2/blob/master/docs/tutorial.md>
- Using R with Hadoop
 - <http://www.revolutionanalytics.com/free-webinars/using-r-hadoop>

Apache Spark™

Lightning-Fast Cluster Computing

61/66

技術準備

- 不怕接觸 Linux (Ubuntu/ Centos)
- 有程式設計開發經驗(C#, Java, Python, SQL, R)
- 基本Java或Scala的概念
- 熟悉 SQL 語法與應用。

Apache Spark provides Resilient Distributed Datasets (RDDs) that store data in memory across a Hadoop cluster. This allows data to be read from HDFS once and to be used many times in order to dramatically improve the performance of interactive tasks like data exploration and iterative algorithms like gradient descent or k-means clustering.

Spark在記憶體內執行程式的運算速度能做到比Hadoop MapReduce的運算速度快上100倍，即便是執行程式於硬碟時，Spark也能快上10倍速度。

[https://zh.wikipedia.org/wiki/Apache_Spark]

The screenshot shows the Apache Spark website with the following content:

- Apache Spark™** Lightning-fast cluster computing
- Navigation: Download, Libraries, Documentation, Examples, Community, FAQ, Apache Software Foundation
- Apache Spark™** is a fast and general engine for large-scale data processing.
- Speed**: Run programs up to 100x faster than Hadoop MapReduce in memory, or 10x faster on disk. A bar chart shows Hadoop at 110s and Spark at 0.9s.
- Ease of Use**: Write applications quickly in Java, Scala, Python, R. Spark offers over 80 high-level operators.
- Generality**: Combine SQL, streaming, and complex analytics. Spark powers a stack of libraries including SQL and DataFrames.
- Latest News**: Preview release of Spark 2.0 (May 26, 2016), Spark Summit (June 6, 2016, San Francisco) agenda posted (Apr 17, 2016), Spark 1.6.1 released (Mar 09, 2016), Submission is open for Spark Summit San Francisco (Feb 11, 2016).

The document titled "SparkR (R on Spark)" shows the following structure:

- Overview
- SparkR DataFrames
 - Starting Up: SparkContext, SQLContext
 - Starting Up from RStudio
 - Creating DataFrames
 - From local data frames
 - From Data Sources
 - From Hive tables
 - DataFrame Operations

講義下載

<http://www.hmwu.idv.tw/>

- 要學的東西太多，但時間太少。
- 「程式可以跑就好啦！幹嘛要寫的好看、寫的有效率啊？」
 - 心態問題：
事情有做就好 vs 把事情做好。
- 統計(角色、任務等等)在哪裡？
- 「如果你本來就認為數據有大用
那剩下的就是技術問題；如果你
本來就認為數據無大用，那剩下的
的都是問題」 (精誠集團Etu負責人蔣居裕)

「語文永遠不是阻止我們了解事物的最大障礙，求知欲才是！」



練習1: 不同R函式計算n!所需的時間^{63/66}

```
factorial.for <- function(n){  
  f <- 1  
  if(n<2) return(1)  
  for(i in 2:n){  
    f <- f*i  
  }  
  f  
}  
factorial.for(5)
```

```
factorial.while <- function(n){  
  f <- 1  
  t <- n  
  while(t>1){  
    f <- f*t  
    t <- t-1  
  }  
  return(f)  
}  
factorial.while(5)
```

```
factroial.repeat <- function(n){  
  f <- 1  
  t <- n  
  repeat{  
    if(t<2) break  
    f <- f*t  
    t <- t-1  
  }  
  return(f)  
}  
factroial.repeat(5)
```

```
factorial.call <- function(n, f){  
  
  if(n <= 1){  
    return(f)  
  }  
  else{  
    factorial.call(n-1, n*f)  
  }  
}  
factorial.call(5, 1)
```

```
factorial.cumprod <- function(n){  
  max(cumprod(1:n))  
}  
factorial.cumprod(5)
```

```
my.factorial <- function(n){  
  ...  
}  
Rprof("output.txt")  
ans <- my.factorial(5000)  
Rprof(NULL)  
summaryRprof("output.txt")
```



練習2

64/66

- 僅輸入所需要的部份資料，而不是全部。

```
Variables <- c("NULL", "NULL", "factor", "numeric")  
myData <- read.table("fileName", colClasses = Variables)
```

- 用適合的函式或演算法: $O(N)$ vs $O(N^2)$

```
x <- 1:10000; s <- sample(x, 10)  
a1 <- which(x %in% s)  
a2 <- intersect(x, s)  
a3 <- which(is.element(x, s))  
  
for(i in 1:10000){  
  for(j in 1:10){  
    if(all.equal(x[i], s[j])){  
      ...  
    }  
  }  
}
```

```
> n <- 10000  
> p <- 1000  
> Mat <- matrix(rnorm(n*p),  
  nrow=n, ncol=p)  
> system.time(apply(Mat, 1, sum))  
  user  system elapsed  
 0.61    0.19    2.56  
> system.time(rowSums(Mat))  
  user  system elapsed  
 0.05    0.00    0.08
```



練習3

65/66

- 資料儲存以二進位檔(binary)為優先:
 - 讀寫文字檔比壓縮二進位檔慢。
 - 壓縮二進位檔又比二進位慢。

```
> n <- 1000  
> p <- 1000  
> Mat <- matrix(rnorm(n*p),  
nrow=n, ncol=p)
```

```
> system.time(write.table(Mat, file="myData.txt"))  
user system elapsed  
8.89 0.09 12.14  
> system.time(read.table("myData.txt"))  
user system elapsed  
10.85 0.06 11.98
```

```
> system.time(save(Mat, file="myData.gz"))  
user system elapsed  
1.11 0.01 2.52  
> system.time(load("myData.gz"))  
user system elapsed  
0.36 0.02 3.56
```

```
> system.time(save(Mat, file="myData.Rdata", compress=FALSE))  
user system elapsed  
0.24 0.00 0.23  
> system.time(load("myData.Rdata"))  
user system elapsed  
0.23 0.00 0.24
```

- 初始向量變數時就先給定長度。避免大量回圈(loop)。
- 採用向量化的運算方式。使用Rprof()檢查程式區塊。

```
comp1 <- function(n){
  x <- numeric()
  for(i in 1:n){
    x[i] <- log(i);
  }
}
```

```
comp2 <- function(n){
  x <- numeric(n)
  for(i in 1:n){
    x[i] <- log(i);
  }
}
```

```
comp3 <- function(n){
  x <- numeric(n)
  for(i in 1:n) {
    x[i] <- log(i);
  }
}
```

```
comp4 <- function(n){
  i <- 1:n
  x <- log(i)
}
```

```
mycomp <- function(n){
  comp1(n)
  comp2(n)
  comp3(n)
  comp4(n)
}
```

```
> Rprof()
> mycomp(10000)
> Rprof(NULL)
> print(summaryRprof())
$by.self
      self.time self.pct total.time total.pct
comp1         0.46   79.3         0.46   79.3
comp2         0.06   10.3         0.06   10.3
comp3         0.06   10.3         0.06   10.3
mycomp        0.00    0.0         0.58  100.0

$by.total
      total.time total.pct self.time self.pct
mycomp         0.58   100.0         0.00    0.0
comp1          0.46   79.3         0.46   79.3
comp2          0.06   10.3         0.06   10.3
comp3          0.06   10.3         0.06   10.3

$sampling.time
[1] 0.58
```